


exampl1

September 14, 2023

```
[125]: from sklearn.datasets import make_classification
import seaborn as sns
import numpy as np
# Generate Clean data
X,y = make_classification(n_samples=500, n_features=2,
    ↪n_informative=2,n_redundant=0, n_repeated=0, n_classes=2,
    ↪n_clusters_per_class=1,class_sep=2,flip_y=0,weights=[0.5,0.5],
    ↪random_state=17)#

sns.scatterplot(x=X[:,0],y=X[:,1],hue=y);
```



```
[126]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X,
    y,
```

```

test_size=0.2, # 20% test,
↳80% train
random_state=42) # make the
↳random split reproducible

len(X_train), len(X_test), len(y_train), len(y_test)

```

[126]: (400, 100, 400, 100)

[24]:

```

[127]: import torch
X_train = torch.from_numpy(X_train).type(torch.float)
y_train = torch.from_numpy(y_train).type(torch.float)
X_test = torch.from_numpy(X_test).type(torch.float)
y_test = torch.from_numpy(y_test).type(torch.float)

```

```

[128]: #Building a model
import torch
from torch import nn

# Make device agnostic code
device = "cuda" if torch.cuda.is_available() else "cpu"
device

```

[128]: 'cpu'

```

[138]: class myModel(nn.Module):
    def __init__(self):
        super().__init__()
        # 2. Create 2 nn.Linear layers capable of handling X and y input and
        ↳output shapes
        self.layer_1 = nn.Linear(in_features=2, out_features=4) # takes in 2
        ↳features (X), produces 5 features
        self.layer_2 = nn.Linear(in_features=4, out_features=1) # takes in 5
        ↳features, produces 1 feature (y)

        # 3. Define a forward method containing the forward pass computation
        def forward(self, x):
            # Return the output of layer_2, a single feature, the same shape as y
            return self.layer_2(self.layer_1(x)) # computation goes through layer_1
            ↳first then the output of layer_1 goes through layer_2

        # 4. Create an instance of the model and send it to target device
        model_0 = myModel().to(device)
        model_0

```

```
[138]: myModel(
        (layer_1): Linear(in_features=2, out_features=4, bias=True)
        (layer_2): Linear(in_features=4, out_features=1, bias=True)
    )
```

```
[139]: state_dict = model_0.state_dict()
        print(state_dict)

OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
        0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
        -0.0706]])), ('layer_2.bias', tensor([0.3854]))])
```

```
[131]: #Setup loss function and optimizer
        # Create a loss function
        # loss_fn = nn.BCELoss() # BCELoss = no sigmoid built-in
        loss_fn = nn.BCEWithLogitsLoss() # BCEWithLogitsLoss = sigmoid built-in

        # Create an optimizer
        optimizer = torch.optim.SGD(params=model_0.parameters(),
                                    lr=0.1)
```

```
[ ]:
```

```
[132]: # Calculate accuracy (a classification metric)
        def accuracy_fn(y_true, y_pred):
            correct = torch.eq(y_true, y_pred).sum().item() # torch.eq() calculates
            ↪ where two tensors are equal
            acc = (correct / len(y_pred)) * 100
            return acc
```

```
[133]: # Train model
        # View the frist 5 outputs of the forward pass on the test data
        y_logits = model_0(X_test.to(device))[:5]
        y_logits
```

```
[133]: tensor([[ 1.1118],
        [ 1.1551],
        [ 0.2263],
        [-0.1948],
        [-0.0858]], grad_fn=<SliceBackward0>)
```

```
[134]: # Use sigmoid on model logits
        y_pred_probs = torch.sigmoid(y_logits)
        y_pred_probs
```

```
[134]: tensor([[0.7525],
              [0.7605],
              [0.5563],
              [0.4515],
              [0.4786]], grad_fn=<SigmoidBackward0>)
```

```
[162]: # Find the predicted labels (round the prediction probabilities)
y_preds = torch.round(y_pred_probs)

# In full
y_pred_labels = torch.round(torch.sigmoid(model_0(X_test.to(device))[:5]))

# Check for equality
print(torch.eq(y_preds.squeeze(), y_pred_labels.squeeze()))

# Get rid of extra dimension
y_preds.squeeze()
```

```
tensor([ True,  True, False, False, False])
```

```
[162]: tensor([1., 1., 1., 0., 0.], grad_fn=<SqueezeBackward0>)
```

```
[163]: import torch
torch.manual_seed(42)
epochs = 5

# Put data to target device

# Build training and evaluation loop
for epoch in range(epochs):
    ### Training
    model_0.train()

    # 1. Forward pass (model outputs raw logits)
    y_logits = model_0(X_train).squeeze() # squeeze to remove extra `1`
    ↪ dimensions, this won't work unless model and data are on same device
    y_pred = torch.round(torch.sigmoid(y_logits)) # turn logits -> pred probs
    ↪ -> pred labls

    # 2. Calculate loss/accuracy
    # loss = loss_fn(torch.sigmoid(y_logits), # Using nn.BCELoss you need torch.
    ↪ sigmoid())
    #
    # y_train)
    loss = loss_fn(y_logits, # Using nn.BCEWithLogitsLoss works with raw logits
    y_train)
    #print(y_logits)
    acc = accuracy_fn(y_true=y_train,
    y_pred=y_pred)
```

```

# 3. Optimizer zero grad
optimizer.zero_grad()

# 4. Loss backwards
loss.backward()

# 5. Optimizer step
optimizer.step()
### Testing
model_0.eval()
with torch.inference_mode():
    # 1. Forward pass
    test_logits = model_0(X_test).squeeze()
    test_pred = torch.round(torch.sigmoid(test_logits))
    # 2. Caculate loss/accuracy
    test_loss = loss_fn(test_logits,
                        y_test)
    test_acc = accuracy_fn(y_true=y_test,
                        y_pred=test_pred)

# Print out what's happening every epochs
print(f"Epoch: {epoch} | Loss: {loss:.5f}, Accuracy: {acc:.2f}% | Test loss:
↪ {test_loss:.5f}, Test acc: {test_acc:.2f}%")
state_dict = model_0.state_dict()
print(state_dict)

```

Epoch: 0 | Loss: 0.54902, Accuracy: 86.00% | Test loss: 0.53172, Test acc: 91.00%

```
OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
        0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
        -0.0706]])), ('layer_2.bias', tensor([0.3854]))])
```

Epoch: 1 | Loss: 0.54902, Accuracy: 86.00% | Test loss: 0.53172, Test acc: 91.00%

```
OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
        0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
        -0.0706]])), ('layer_2.bias', tensor([0.3854]))])
```

Epoch: 2 | Loss: 0.54902, Accuracy: 86.00% | Test loss: 0.53172, Test acc: 91.00%

```
OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
```

```

        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
-0.0706]])), ('layer_2.bias', tensor([0.3854]))))
Epoch: 3 | Loss: 0.54902, Accuracy: 86.00% | Test loss: 0.53172, Test acc:
91.00%
OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
-0.0706]])), ('layer_2.bias', tensor([0.3854]))))
Epoch: 4 | Loss: 0.54902, Accuracy: 86.00% | Test loss: 0.53172, Test acc:
91.00%
OrderedDict([('layer_1.weight', tensor([[ 0.5406,  0.5869],
        [-0.1657,  0.6496],
        [-0.1549,  0.1427],
        [-0.3443,  0.4153]])), ('layer_1.bias', tensor([ 0.6233, -0.5188,
0.6146,  0.1323])), ('layer_2.weight', tensor([[ 0.3694,  0.0677,  0.2411,
-0.0706]])), ('layer_2.bias', tensor([0.3854]))))

```

```

[101]: loss = nn.CrossEntropyLoss()
        output = loss(input, target)
        output.backward()

```

```

[73]: print(input, '\n', target, '\n', output)

```

```

tensor([[ 0.3367,  0.1288,  0.2345,  0.2303, -1.1229],
        [-0.1863,  2.2082, -0.6380,  0.4617,  0.2674],
        [ 0.5349,  0.8094,  1.1103, -1.6898, -0.9890]], requires_grad=True)
tensor([0, 4, 3])
tensor(2.4607, grad_fn=<NllLossBackward0>)

```