

linear_ridge_lasso_elasticNet_regression

March 30, 2024

```
[1]: ## Aditya Agre
    ## 121B1B006
    ## second part: Linear Regression, Ridge, Lasso, and ElasticNet models
    ## Diabetes dataset

[2]: import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt
    from sklearn import datasets

[3]: diabetes = datasets.load_diabetes()
    diabetes

[3]: {'data': array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
    0.01990749, -0.01764613],
    [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
    -0.06833155, -0.09220405],
    [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
    0.00286131, -0.02593034],
    ...,
    [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
    -0.04688253,  0.01549073],
    [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
    0.04452873, -0.02593034],
    [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
    -0.00422151,  0.00306441]]),
    'target': array([151.,  75., 141., 206., 135.,  97., 138.,  63., 110., 310.,
    101.,
    69., 179., 185., 118., 171., 166., 144.,  97., 168.,  68.,  49.,
    68., 245., 184., 202., 137.,  85., 131., 283., 129.,  59., 341.,
    87.,  65., 102., 265., 276., 252.,  90., 100.,  55.,  61.,  92.,
    259.,  53., 190., 142.,  75., 142., 155., 225.,  59., 104., 182.,
    128.,  52.,  37., 170., 170.,  61., 144.,  52., 128.,  71., 163.,
    150.,  97., 160., 178.,  48., 270., 202., 111.,  85.,  42., 170.,
    200., 252., 113., 143.,  51.,  52., 210.,  65., 141.,  55., 134.,
    42., 111.,  98., 164.,  48.,  96.,  90., 162., 150., 279.,  92.,
    83., 128., 102., 302., 198.,  95.,  53., 134., 144., 232.,  81.,
    104.,  59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,
```

```

173., 180., 84., 121., 161., 99., 109., 115., 268., 274., 158.,
107., 83., 103., 272., 85., 280., 336., 281., 118., 317., 235.,
60., 174., 259., 178., 128., 96., 126., 288., 88., 292., 71.,
197., 186., 25., 84., 96., 195., 53., 217., 172., 131., 214.,
59., 70., 220., 268., 152., 47., 74., 295., 101., 151., 127.,
237., 225., 81., 151., 107., 64., 138., 185., 265., 101., 137.,
143., 141., 79., 292., 178., 91., 116., 86., 122., 72., 129.,
142., 90., 158., 39., 196., 222., 277., 99., 196., 202., 155.,
77., 191., 70., 73., 49., 65., 263., 248., 296., 214., 185.,
78., 93., 252., 150., 77., 208., 77., 108., 160., 53., 220.,
154., 259., 90., 246., 124., 67., 72., 257., 262., 275., 177.,
71., 47., 187., 125., 78., 51., 258., 215., 303., 243., 91.,
150., 310., 153., 346., 63., 89., 50., 39., 103., 308., 116.,
145., 74., 45., 115., 264., 87., 202., 127., 182., 241., 66.,
94., 283., 64., 102., 200., 265., 94., 230., 181., 156., 233.,
60., 219., 80., 68., 332., 248., 84., 200., 55., 85., 89.,
31., 129., 83., 275., 65., 198., 236., 253., 124., 44., 172.,
114., 142., 109., 180., 144., 163., 147., 97., 220., 190., 109.,
191., 122., 230., 242., 248., 249., 192., 131., 237., 78., 135.,
244., 199., 270., 164., 72., 96., 306., 91., 214., 95., 216.,
263., 178., 113., 200., 139., 139., 88., 148., 88., 243., 71.,
77., 109., 272., 60., 54., 221., 90., 311., 281., 182., 321.,
58., 262., 206., 233., 242., 123., 167., 63., 197., 71., 168.,
140., 217., 121., 235., 245., 40., 52., 104., 132., 88., 69.,
219., 72., 201., 110., 51., 277., 63., 118., 69., 273., 258.,
43., 198., 242., 232., 175., 93., 168., 275., 293., 281., 72.,
140., 189., 181., 209., 136., 261., 113., 131., 174., 257., 55.,
84., 42., 146., 212., 233., 91., 111., 152., 120., 67., 310.,
94., 183., 66., 173., 72., 49., 64., 48., 178., 104., 132.,
220., 57.]),

```

```
'frame': None,
```

```
'DESCR': '.._diabetes_dataset:\n\nDiabetes dataset\n-----\n\nTen
baseline variables, age, sex, body mass index, average blood\npressure, and six
blood serum measurements were obtained for each of n =\n442 diabetes patients,
as well as the response of interest, a\nquantitative measure of disease
progression one year after baseline.\n\n**Data Set Characteristics:**\n\n
: Number of Instances: 442\n\n : Number of Attributes: First 10 columns are
numeric predictive values\n\n : Target: Column 11 is a quantitative measure of
disease progression one year after baseline\n\n : Attribute Information:\n
- age      age in years\n      - sex\n      - bmi      body mass index\n      - bp
average blood pressure\n      - s1      tc, total serum cholesterol\n      - s2
ldl, low-density lipoproteins\n      - s3      hdl, high-density lipoproteins\n
- s4      tch, total cholesterol / HDL\n      - s5      ltcg, possibly log of
serum triglycerides level\n      - s6      glu, blood sugar level\n\nNote: Each
of these 10 feature variables have been mean centered and scaled by the standard
deviation times the square root of `n_samples` (i.e. the sum of squares of each
column totals 1).\n\nSource
```

```

URL:\nhttps://www4.stat.ncsu.edu/~boos/var.select/diabetes.html\n\nFor more
information see:\nBradley Efron, Trevor Hastie, Iain Johnstone and Robert
Tibshirani (2004) "Least Angle Regression," Annals of Statistics (with
discussion),
407-499.\n(https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf)\n',
'feature_names': ['age',
'sex',
'bmi',
'bp',
's1',
's2',
's3',
's4',
's5',
's6'],
'data_filename': 'diabetes_data_raw.csv.gz',
'target_filename': 'diabetes_target.csv.gz',
'data_module': 'sklearn.datasets.data'}

```

```
[4]: diabetes.feature_names
```

```
[4]: ['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']
```

```
[5]: x = diabetes.data
y = diabetes.target
```

```
[6]: x.shape, y.shape
```

```
[6]: ((442, 10), (442,))
```

```
[7]: from sklearn.model_selection import train_test_split
```

```
[8]: x1, x2, y1, y2 = train_test_split(x,y,test_size=0.3,random_state=99)
```

```
[ ]:
```

```
[9]: # Linear Regression
from sklearn.linear_model import LinearRegression
```

```
[10]: lr_object = LinearRegression()
```

```
[11]: lr_object.fit(x1,y1)
```

```
[11]: LinearRegression()
```

```
[12]: predictions_on_x2 = lr_object.predict(x2)
```

```
[13]: comparison_Linear_regression = pd.DataFrame({'Given_values_y2' : y2,
↳ 'predicted_values' : predictions_on_x2})
print(comparison_Linear_regression)
```

	Given_values_y2	predicted_values
0	75.0	77.999103
1	128.0	170.447121
2	125.0	109.036606
3	332.0	223.843071
4	37.0	87.384304
..
128	48.0	202.349391
129	172.0	144.663523
130	51.0	82.400892
131	277.0	185.443544
132	94.0	101.563213

[133 rows x 2 columns]

```
[27]: lr_object.score(x2,y2)
```

```
[27]: 0.4545709909725648
```

```
[ ]:
```

```
[14]: # Ridge regression. (Still linear)
from sklearn.linear_model import Ridge
```

```
[15]: ridge_obj = Ridge(alpha=0.1)
ridge_obj.fit(x1,y1)
Ridge()
```

```
[15]: Ridge()
```

```
[16]: predictions_on_x2 = ridge_obj.predict(x2)
```

```
[17]: comparison_Ridge_regression = pd.DataFrame({'Given_values_y2' : y2,
↳ 'predicted_values' : predictions_on_x2})
print(comparison_Ridge_regression)
```

	Given_values_y2	predicted_values
0	75.0	89.451853
1	128.0	167.891164
2	125.0	105.681527
3	332.0	219.512632
4	37.0	86.868393
..
128	48.0	192.038147
129	172.0	143.200097

130	51.0	80.603296
131	277.0	184.413866
132	94.0	105.191343

[133 rows x 2 columns]

```
[28]: ridge_obj.score(x2,y2)
```

```
[28]: 0.46062181702662
```

```
[ ]:
```

```
[18]: ## Lasso
      from sklearn.linear_model import Lasso
```

```
[19]: lasso_obj = Lasso(alpha = 0.1)
      lasso_obj.fit(x1,y1)
      Lasso()
```

```
[19]: Lasso()
```

```
[20]: predictions = lasso_obj.predict(x2)
```

```
[21]: comparison_Lasso_regr = pd.DataFrame({'Given_values_y2' : y2,
      ↪ 'predicted_values' : predictions_on_x2})
      print(comparison_Lasso_regr)
```

	Given_values_y2	predicted_values
0	75.0	89.451853
1	128.0	167.891164
2	125.0	105.681527
3	332.0	219.512632
4	37.0	86.868393
..
128	48.0	192.038147
129	172.0	143.200097
130	51.0	80.603296
131	277.0	184.413866
132	94.0	105.191343

[133 rows x 2 columns]

```
[29]: lasso_obj.score(x2,y2)
```

```
[29]: 0.4672409597780036
```

```
[ ]:
```

```
[22]: ## Elastic Net  
from sklearn.linear_model import ElasticNet
```

```
[43]: elastic_object = ElasticNet(alpha = 0.00005)  
elastic_object.fit(x1,y1)  
ElasticNet()
```

```
[43]: ElasticNet()
```

```
[44]: predictions = elastic_object.predict(x2)
```

```
[45]: comparison_elastic_regr = pd.DataFrame({'Given_values_y2' : y2,  
↳ 'predicted_values' : predictions_on_x2})  
print(comparison_elastic_regr)
```

	Given_values_y2	predicted_values
0	75.0	89.451853
1	128.0	167.891164
2	125.0	105.681527
3	332.0	219.512632
4	37.0	86.868393
..
128	48.0	192.038147
129	172.0	143.200097
130	51.0	80.603296
131	277.0	184.413866
132	94.0	105.191343

[133 rows x 2 columns]

```
[46]: elastic_object.score(x2,y2)
```

```
[46]: 0.45312270953163125
```

```
[ ]: ## Therefore, after comparison using R squared test, the best performance is  
↳ shown by Lasso regression.
```